

B.M.S. COLLEGE OF ENGINEERING

(Autonomous College under VTU, Approved by AICTE, Accredited by NAAC)

MASTER OF COMPUTER APPLICATIONS

(Accredited by NBA for 5 years 2019 - 2024)



BIG DATA ANALYTICS (22MCA2PEBD)

LAB REPORT

SUBMITTED BY

Manohar T H

(1BM23MC051)

UNDER THE GUIDANCE OF
S SHILPA

(Assistant Professor)

B.M.S. COLLEGE OF ENGINEERING

(Autonomous College under VTU, Approved by AICTE, Accredited by NAAC)

MASTER OF COMPUTER APPLICATIONS

(Accredited by NBA for 5 years 2019 - 2024)



LABORATORY CERTIFICATE

This is to certify that **MANOHAR T H (1BM23MC051)** has satisfactorily completed the course of practical in “**Big Data Analytics–22MCA2PEBD**” Laboratory prescribed by **BMS College of Engineering** (Autonomous college under VTU) 2nd Semester MCA course in this college during the year 2022 - 2023.

Signature of Batch in charge

S Shilpa
Examiner:

Signature of HOD

Dr. Ch. Ram Mohan Reddy

CONTENTS

SL. No.	Programs	Page No.
1.	Demonstration and installation of HADOOP cluster	1-6
2.	Execution of HDFS Commands for interaction with Hadoop Environment	7-10
3.	Create and execute map reduce programs	11-13
4.	Data Processing Using Hive	14-26
5.	Data processing using Spark	27-32
6.	Programming in Cassandra	33-41

1. Demonstration and installation of HADOOP cluster

Sudo apt update

Step 1: Install Java Development Kit

1. sudo apt update && sudo apt install openjdk-11-jdk
2. java -version
3. dirname \$(dirname \$(readlink -f \$(which java)))
4. sudo adduser hadoop
5. su - hadoop
6. ssh-keygen -t rsa
7. cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
8. chmod 640 ~/.ssh/authorized_keys
9. sudo adduser hadoop sudo
10. sudo apt install openssh-server
11. ssh localhost
12. wget https://dlcdn.apache.org/hadoop/common/hadoop-3.3.4/hadoop-3.3.4.tar.gz
13. tar xzf hadoop-3.3.4.tar.gz
14. mv hadoop-3.3.4 hadoop

nano ~/.bashrc

```
export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64      export
HADOOP_HOME=/home/hadoop/hadoop                          export
HADOOP_INSTALL=$HADOOP_HOME                              export
HADOOP_MAPRED_HOME=$HADOOP_HOME                          export
HADOOP_COMMON_HOME=$HADOOP_HOME                          export
HADOOP_HDFS_HOME=$HADOOP_HOME                            export
HADOOP_YARN_HOME=$HADOOP_HOME                            export
HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native    export
PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin           export
HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib/native"
```

source ~/.bashrc nano \$HADOOP_HOME/etc/hadoop/hadoop-env.sh

```
export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64
```

Step 2: Configuring Hadoop

```
mkdir -p ~/hadoopdata/hdfs/{namenode,datanode}
```

1. nano \$SHADOOP_HOME/etc/hadoop/core-site.xml

```
configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

2. nano \$SHADOOP_HOME/etc/hadoop/hdfs-site.xml

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.name.dir</name>
    <value>file:///home/hadoop/hadoopdata/hdfs/namenode</value>
  </property>
  <property>
    <name>dfs.data.dir</name>
    <value>file:///home/hadoop/hadoopdata/hdfs/datanode</value>
  </property>
</configuration>
```

3. nano \$SHADOOP_HOME/etc/hadoop/mapred-site.xml

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

4. nano \$SHADOOP_HOME/etc/hadoop/yarn-site.xml

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>
```

Step 3: Start Hadoop Cluster

1.hdfs namenode -

format start-all.sh

<http://localhost:9870>

<http://localhost:8088>

1 Execution of HDFS Commands for interaction with Hadoop Environment

1. Create a directory

```
hadoop fs -mkdir /manu
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop$ hadoop fs -ls / manu/
Found 4 items
drwxr-xr-x - hadoop supergroup      0 2024-07-03 15:54 /dir2
drwxr-xr-x - hadoop supergroup      0 2024-07-09 11:49 /manu
drwxr-xr-x - hadoop supergroup      0 2024-07-09 11:34 /manu1
drwxr-xr-x - hadoop supergroup      0 2024-07-03 15:49 /newDir
ls: 'manu/': No such file or directory
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop$ hadoop fs -ls /manu/
Found 1 items
-rw-r--r-- 1 hadoop supergroup      0 2024-07-09 11:49 /manu/empty.txt
```

2. Create an empty file

```
$~ hadoop fs -touchz /manu/empty.txt
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop$ hadoop fs -ls / manu/
Found 4 items
drwxr-xr-x - hadoop supergroup      0 2024-07-03 15:54 /dir2
drwxr-xr-x - hadoop supergroup      0 2024-07-09 11:49 /manu
drwxr-xr-x - hadoop supergroup      0 2024-07-09 11:34 /manu1
drwxr-xr-x - hadoop supergroup      0 2024-07-03 15:49 /newDir
ls: 'manu/': No such file or directory
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop$ hadoop fs -ls /manu/
Found 1 items
-rw-r--r-- 1 hadoop supergroup      0 2024-07-09 11:49 /manu/empty.txt
```

3. List all the files in a directory, recursively displays entries in all subdirectories of a path

```
hadoop fs -ls /manu/
```

```
hadoop fs -ls /
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop$ hadoop fs -ls / manu/
Found 4 items
drwxr-xr-x - hadoop supergroup      0 2024-07-03 15:54 /dir2
drwxr-xr-x - hadoop supergroup      0 2024-07-09 11:49 /manu
drwxr-xr-x - hadoop supergroup      0 2024-07-09 11:34 /manu1
drwxr-xr-x - hadoop supergroup      0 2024-07-03 15:49 /newDir
ls: 'manu/': No such file or directory
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop$ hadoop fs -ls /manu/
Found 1 items
-rw-r--r-- 1 hadoop supergroup      0 2024-07-09 11:49 /manu/empty.txt
```

4. Copy files/folders from local file system to hdfs store

```
hadoop fs -put /home/hadoop/new.txt /manu/  
hadoop fs -ls /manu/
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /manu/  
Found 2 items  
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 11:49 /manu/empty.txt  
-rw-r--r-- 1 hadoop supergroup 28 2024-07-09 12:07 /manu/new.txt  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$
```

5. Print the file contents

```
hadoop fs -cat /manu/new.txt
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /manu/new.txt  
hello  
this is local file
```

6. To copy files/folders from hdfs store to local file system

```
hadoop fs -get /manu/hdfs-file.txt /home/hadoop
```

```
Desktop Documents Downloads hadoop hadoop-3.3.6.tar.gz hadoopdata manu Music new.txt Pictures Public snap Templates Videos  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -get /manu/hdfs-file.txt /home/hadoop  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ ls  
Desktop Documents Downloads hadoop hadoop-3.3.6.tar.gz hadoopdata hdfs-file.txt manu Music new.txt Pictures Public snap Templates Videos  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$
```

7. Move file from local to hdfs

```
hadoop fs -moveFromLocal /home/hadoop/new.txt /manu/
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -moveFromLocal /home/hadoop/new.txt /manu/  
moveFromLocal: `/manu/new.txt': File exists  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -rm /manu/new.txt  
Deleted /manu/new.txt  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -moveFromLocal /home/hadoop/new.txt /manu/  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /manu/  
Found 3 items  
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 11:49 /manu/empty.txt  
-rw-r--r-- 1 hadoop supergroup 44 2024-07-09 12:26 /manu/hdfs-file.txt  
-rw-r--r-- 1 hadoop supergroup 28 2024-07-09 12:44 /manu/new.txt  
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$
```


8. Copy files within hdfs

`hadoop fs -moveFromLocal /home/hadoop/new.txt /manu/`

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -touchz /manu/cpoied.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cp /manu/hdfs-file.txt /manu/copied.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /manu/copied.txt
helo
this is hdfs file
in hadoop directory
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /manu/hdfs-file.txt
helo
this is hdfs file
in hadoop directory
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$
```

9. Move/rename files within hdfs

`hadoop fs -mv /manu/cpoied.txt /manu/movied.txt`

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /manu/
Found 3 items
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 12:48 /manu/cpoied.txt
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 11:49 /manu/empty.txt
-rw-r--r-- 1 hadoop supergroup 28 2024-07-09 12:44 /manu/new.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -mv /manu/cpoied.txt /manu/movied.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /manu/
Found 3 items
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 11:49 /manu/empty.txt
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 12:48 /manu/movied.txt
-rw-r--r-- 1 hadoop supergroup 28 2024-07-09 12:44 /manu/new.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$
```

10. Delete a file, delete a file from HDFS recursively

`hadoop fs -rmr /manu/new.txt`

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /manu/
Found 3 items
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 11:49 /manu/empty.txt
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 12:48 /manu/movied.txt
-rw-r--r-- 1 hadoop supergroup 28 2024-07-09 12:44 /manu/new.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -rmr /manu/new.txt
rmr: DEPRECATED: Please use '-rm -r' instead.
Deleted /manu/new.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /manu/
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 11:49 /manu/empty.txt
-rw-r--r-- 1 hadoop supergroup 0 2024-07-09 12:48 /manu/movied.txt
```

11. Display Size of directory/file, size of each file in directory

It will give the size of each file in directory.

```
hadoop fs -du /
```

This command will give the total size of directory/file.

```
hadoop fs -dus /manu/
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -du /manu/
0 0 /manu/empty.txt
0 0 /manu/movied.txt
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -dus /manu/
dus: DEPRECATED: Please use 'du -s' instead.
0 0 /manu
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -du /
0 0 /dir2
0 0 /manu
0 0 /manu.empty.txt
0 0 /manu1
0 0 /newDir
0 0 /user
```

12. Append a file in hdfs

```
$~ hadoop fs -appendToFile /manu/new.txt home//hadoop/data.txt
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /manu/new.txt
heloo
this is local file
```

3. Create and execute map reduce programs

WC_Mapper.java

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
public class WC_Mapper extends MapReduceBase implements
Mapper<LongWritable,Text,Text,IntWritable>{
private final static IntWritable one = new IntWritable(1);
private Text word = new Text();
public void map(LongWritable key, Text value,OutputCollector<Text,IntWritable>
output,
Reporter reporter) throws IOException{
String line = value.toString();
StringTokenizer tokenizer = new StringTokenizer(line);
while (tokenizer.hasMoreTokens()){
word.set(tokenizer.nextToken());
output.collect(word, one);
}
}}
```

WC_Reducer.java

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
public class WC_Reducer extends MapReduceBase implements
Reducer<Text,IntWritable,Text,IntWritable> {
public void reduce(Text key,
Iterator<IntWritable>
values,OutputCollector<Text,IntWritable> output,
Reporter reporter) throws IOException {
int sum=0;
while (values.hasNext()) {
```

```
sum+=values.next().get();
}
output.collect(key,new IntWritable(sum));
}
}
```

WC_Runner.java

```
import java.io.IOException;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;
public class WC_Runner {
public static void main(String[] args) throws IOException{
JobConf conf = new JobConf(WC_Runner.class);
conf.setJobName("WordCount");
conf.setOutputKeyClass(Text.class);
conf.setOutputValueClass(IntWritable.class);
conf.setMapperClass(WC_Mapper.class);
conf.setCombinerClass(WC_Reducer.class);
conf.setReducerClass(WC_Reducer.class);
conf.setInputFormat(TextInputFormat.class);
conf.setOutputFormat(TextOutputFormat.class);
FileInputFormat.setInputPaths(conf,new Path(args[0]));
FileOutputFormat.setOutputPath(conf,new Path(args[1]));
JobClient.runJob(conf);
}
}
```

Output:

```
cat > new
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ cat > new
Hello everyone how are you
good morning iam fine
what about you
that is super
tahnk you
```

```
hadoop fs -mkdir /lab1
hadoop fs -put /home/hadoop/new /lab1/new
hadoop fs -cat /lab1/new
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -mkdir /lab1
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -put /home/hadoop/new /lab1/new
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /lab1/new
Hello everyone how are you
good morning iam fine
what about you
that is super
tahnk you
```

```
hadoop jar /home/hadoop/hadoop/WC_Mapreduce.jar WC_Runner /lab1/new
/WC_output1
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop jar /home/hadoop/hadoop/WC_Mapreduce.jar WC_Runner /lab1/new /WC_output1
2024-07-30 12:49:30,927 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2024-07-30 12:49:30,967 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2024-07-30 12:49:30,967 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2024-07-30 12:49:30,974 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
```

```
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=90
File Output Format Counters
  Bytes Written=110
```

```
hadoop fs -ls WOutput
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls WOutput
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2024-07-30 12:48 WOutput/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 25 2024-07-30 12:48 WOutput/part-000000
```

```
hadoop fs -cat /WOutput1/part-00000
```

```
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/manu$ hadoop dfs -cat WOutput/part-00000
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.
am      1
helooo  1
i       1
manu    1
hadoop@mca-HP-Elite-Tower-800-G9-Desktop-PC:~/manu$
```

4. Data Processing Using Hive

Database name: Cabbase

1. Cabbase

```
hive> create database if not exists Cab_base;
OK
Time taken: 0.084 seconds
hive> show databases;
OK
cab_base
default
Time taken: 0.015 seconds, Fetched: 2 row(s)
hive> use cab_base;
OK
Time taken: 0.013 seconds
```

cab

<u>cab_id</u>	cab_number	cab_type	cab_status	cab_price
---------------	------------	----------	------------	-----------

```
hive> create table if not exists cab(
  > cab_id int,
  > cab_number int,
  > cab_type string,
  > cab_status string,
  > cab_price int)
  > row format delimited
  > fields terminated by '\t'
  > lines terminated by '\n';
```

```
hive> desc cab;
OK
cab_id          int
cab_number      int
cab_type        string
cab_status      string
cab_price       int
Time taken: 0.091 seconds, Fetched: 5 row(s)
```

Inserting Data Into Table :

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/cab.txt' into table
cab;
```

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/cab.txt' into table cab;
Loading data to table cab_base.cab
OK
Time taken: 0.891 seconds
```

1. Display the contents

hive> select * from cab;

```
hive> select * from cab;
OK
101      123456  Sedan   Available      500
102      987654  SUV     Available      800
103      456789  Minivan Unavailable    600
104      159357  Hatchback Available      400
105      753951  Sedan   Unavailable    550
106      951357  SUV     Available      850
107      852741  Minivan Available      650
108      147369  Hatchback Unavailable    450
109      963852  Sedan   Available      600
110      741852  SUV     Available      900
Show Applications seconds, Fetched: 10 row(s)
```

Driver

<u>driver_id</u>	driver_name	driver_sal
------------------	-------------	------------

```
hive> create table if not exists Driver(
> driver_id int,
> driver_name string,
> driver_sal int)
> row format delimited
> fields terminated by '\t'
> lines terminated by '\n';
```

```
Time taken: 0.776 seconds, Fetched: 10 row(s)
hive> create table if not exists Driver(
> driver_id int,
> driver_name string,
> driver_sal int)
> row format delimited
> fields terminated by '\t'
> lines terminated by '\n';
OK
Time taken: 0.078 seconds
```

hive> desc Driver;

```
hive> desc Driver;
OK
driver_id      int
driver_name    string
driver_sal     int
Time taken: 0.033 seconds, Fetched: 3 row(s)
hive>
```

hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/Driver.txt' into table Driver;

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/Driver.txt' into table Driver;
Loading data to table cab_base.driver
OK
Time taken: 0.511 seconds
```

1. Display the contents

hive> select * from Driver;

```
Time taken: 0.511 seconds
hive> select * from Driver;
OK
1001    Aarav Sharma    30000
1002    Ananya Gupta    35000
1003    Rohan Patel     32000
1004    Diya Verma       28000
1005    Arjun Mehta      40000
1006    Sneha Reddy      45000
1007    Krish Iyer       37000
1008    Pooja Singh      36000
1009    Aditya Joshi     29000
1010    Neha Nair        31000
Time taken: 0.067 seconds, Fetched: 10 row(s)
```


Rental

<u>rental_id</u>	rental_date	rental_time	rental_dest	Payment
------------------	-------------	-------------	-------------	---------

```
hive> CREATE TABLE IF NOT EXISTS Rental (
  > rental_id INT,
  > rental_date STRING,
  > rental_time STRING,
  > rental_dest STRING,
  > Payment INT )
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY '\t'
  > LINES TERMINATED BY '\n';
```

```
hive> CREATE TABLE IF NOT EXISTS Rental (
  > rental_id INT,
  > rental_date STRING,
  > rental_time STRING,
  > rental_dest STRING,
  > Payment INT
  > )
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY '\t'
  > LINES TERMINATED BY '\n'
  > ;
```

```
hive> desc Rental;
```

```
hive> desc Rental;
OK
rental_id          int
rental_date        string
rental_time        string
rental_dest        string
payment            int
Time taken: 0.438 seconds, Fetched: 5 row(s)
hive>
```

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/rental.txt' into table
Rental;
```

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/rental.txt' into table Rental;
Loading data to table default.rental
OK
Time taken: 0.414 seconds
```

1. Display the contents

hive> select * from Rental;

```
hive> select * from Rental;
OK
1      2024-08-01      10:00 AM      Mumbai 1500
2      2024-08-02      11:30 AM      Delhi 2000
3      2024-08-03      09:15 AM      Bangalore 1800
4      2024-08-04      02:45 PM      Chennai 1600
5      2024-08-05      03:30 PM      Kolkata 1700
6      2024-08-06      12:00 PM      Pune 1900
7      2024-08-07      01:00 PM      Ahmedabad 2100
8      2024-08-08      05:30 PM      Hyderabad 1750
9      2024-08-09      08:00 AM      Surat 1650
10     2024-08-10      04:15 PM      Jaipur 1550
NULL   NULL   NULL   NULL   NULL
Time taken: 0.051 seconds, Fetched: 11 row(s)
```

Customer

<u>cus_id</u>	cus_fname	cus_lname	cus_gender	cus_age	cus_num	cus_email
---------------	-----------	-----------	------------	---------	---------	-----------

hive> CREATE TABLE IF NOT EXISTS Customer (

- > cus_id INT,
- > cus_fname STRING,
- > cus_lname STRING,
- > cus_gender STRING,
- > cus_age INT,
- > cus_num STRING,
- > cus_email STRING)
- > ROW FORMAT DELIMITED
- > FIELDS TERMINATED BY '\t'
- > LINES TERMINATED BY '\n';

```
hive> CREATE TABLE IF NOT EXISTS Customer (
>   cus_id INT,
>   cus_fname STRING,
>   cus_lname STRING,
>   cus_gender STRING,
>   cus_age INT,
>   cus_num STRING,
>   cus_email STRING
> )
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY '\t'
> LINES TERMINATED BY '\n';
OK
Time taken: 0.084 seconds
```

hive> desc Customer;

```
hive> desc Customer;
OK
cus_id          int
cus_fname       string
cus_lname       string
cus_gender      string
cus_age         int
cus_num         string
cus_email       string
Time taken: 0.439 seconds, Fetched: 7 row(s)
```

hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/cost.txt' into table Customer;

1. Display the contents

hive> select * from Customer;

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/cost.txt' into table Customer;
Loading data to table default.customer
OK
Time taken: 0.415 seconds
hive> select * from Customer;
OK
1      Ravi   Sharma  Male   30      9876543210  ravi.sharma@example.com
2      Anita  Patel   Female 28      9876543211  anita.patel@example.com
3      Suresh Kumar  Male   35      9876543212  suresh.kumar@example.com
4      Neha   Verma   Female 26      9876543213  neha.verma@example.com
5      Amit   Singh  Male   40      9876543214  amit.singh@example.com
6      Pooja  Desai   Female 32      9876543215  pooja.desai@example.com
7      Vikram Reddy   Male   29      9876543216  vikram.reddy@example.com
8      Sneha  Mishra  Female 31      9876543217  sneha.mishra@example.com
9      Ajay   Gupta   Male   38      9876543218  ajay.gupta@example.com
10     Sinran Kaur    Female 27      9876543219  sinran.kaur@example.com
Time taken: 0.682 seconds, Fetched: 10 row(s)
```

Transaction

<u>tran_id</u>	tran_name	tran_date	car_id	rental_id	cust_id
----------------	-----------	-----------	--------	-----------	---------

```
hive> CREATE TABLE IF NOT EXISTS Transaction (
  > tran_id INT,
  > tran_name STRING,
  > tran_date STRING,
  > car_id INT,
  > rental_id INT,
  > cust_id INT
  > )
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY '\t'
  > LINES TERMINATED BY '\n';
```

```
hive> desc Transaction;
```

```
hive> CREATE TABLE IF NOT EXISTS Transaction (
  > tran_id INT,
  > tran_name STRING,
  > tran_date STRING,
  > car_id INT,
  > rental_id INT,
  > cust_id INT
  > )
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY '\t'
  > LINES TERMINATED BY '\n';
OK
Time taken: 0.554 seconds
hive> desc Transaction;
OK
tran_id      int
tran_name    string
tran_date    string
car_id       int
rental_id    int
cust_id      int
Time taken: 0.158 seconds, Fetched: 6 row(s)
```

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/trans.txt' into table
Transaction;
```

```
hive> load data local inpath '/home/hadoop/apache-hive-3.1.2-bin/bin/trans.txt' into table Transaction;
Loading data to table default.transaction
OK
Time taken: 0.298 seconds
```

1. Display the contents

hive> select * from Transaction;

```
Time taken: 0.066 seconds
hive> select * from Transaction;
OK
101    Payment Received      2023-04-21      123456    201      301
102    Payment Processed     2024-08-30      987654    202      302
103    Payment Completed     2020-06-13      456789    203      303
104    Payment Confirmed     2019-10-04      159357    204      304
105    Payment Successful     2024-08-05      753951    205      305
106    Payment Verified      2024-02-26      951357    206      306
107    Payment Approved      2024-04-29      852741    207      307
108    Payment Declined       2021-08-08      147369    208      308
109    Payment Reversed       2016-07-09      963852    209      309
110    Payment Refunded       2013-05-10      741852    210      310
Time taken: 0.066 seconds, Fetched: 10 row(s)
```

2. Limit the display to three rows or n no. of rows. cab

hive> SELECT * FROM cab LIMIT 3;

hive> SELECT * FROM cab LIMIT 7;

```
Time taken: 0.067 seconds
hive> SELECT * FROM cab LIMIT 3;
OK
101    123456    Sedan    Available    500
102    987654    SUV      Available    800
103    456789    Minivan  Unavailable  600
Time taken: 0.068 seconds, Fetched: 3 row(s)
hive> SELECT * FROM cab LIMIT 7;
OK
101    123456    Sedan    Available    500
102    987654    SUV      Available    800
103    456789    Minivan  Unavailable  600
104    159357    Hatchback Available    400
105    753951    Sedan    Unavailable  550
106    951357    SUV      Available    850
107    852741    Minivan  Available    650
Time taken: 0.072 seconds, Fetched: 7 row(s)
```

Driver

hive> SELECT * FROM Driver LIMIT 3;

hive> SELECT * FROM Driver LIMIT 8;

```
Time taken: 0.050 seconds
hive> SELECT * FROM Driver LIMIT 3;
OK
1001    Aarav Sharma    30000
1002    Ananya Gupta    35000
1003    Rohan Patel     32000
Time taken: 0.823 seconds, Fetched: 3 row(s)
hive> SELECT * FROM Driver LIMIT 8;
OK
1001    Aarav Sharma    30000
1002    Ananya Gupta    35000
1003    Rohan Patel     32000
1004    Diya Verma      28000
1005    Arjun Mehta     40000
1006    Sneha Reddy     45000
1007    Krish Iyer      37000
1008    Pooja Singh     36000
Time taken: 0.075 seconds, Fetched: 8 row(s)
hive>
```

Rental

hive> SELECT * FROM Rental LIMIT 3;

hive> SELECT * FROM Rental LIMIT 6;

```
hive> SELECT * FROM Rental LIMIT 3;
OK
1      2024-08-01      10:00 AM      Mumbai 1500
2      2024-08-02      11:30 AM      Delhi   2000
3      2024-08-03      09:15 AM      Bangalore 1800
Time taken: 0.086 seconds, Fetched: 3 row(s)
hive> SELECT * FROM Rental LIMIT 6;
OK
1      2024-08-01      10:00 AM      Mumbai 1500
2      2024-08-02      11:30 AM      Delhi   2000
3      2024-08-03      09:15 AM      Bangalore 1800
4      2024-08-04      02:45 PM      Chennai 1600
5      2024-08-05      03:30 PM      Kolkata 1700
6      2024-08-06      12:00 PM      Pune    1900
Time taken: 0.072 seconds, Fetched: 6 row(s)
```

Customer

hive> SELECT * FROM Customer LIMIT 6;

hive> SELECT * FROM Customer LIMIT 3;

```
hive> SELECT * FROM Customer LIMIT 6;
OK
1      Ravi   Sharma  Male   30      9876543210      ravi.sharma@example.com
2      Anita  Patel   Female 28      9876543211      anita.patel@example.com
3      Suresh Kumar Male   35      9876543212      suresh.kumar@example.com
4      Neha   Verma   Female 26      9876543213      neha.verma@example.com
5      Amit   Singh  Male   40      9876543214      amit.singh@example.com
6      Pooja   Desai   Female 32      9876543215      pooja.desai@example.com
Time taken: 0.064 seconds, Fetched: 6 row(s)
hive> SELECT * FROM Customer LIMIT 3;
OK
1      Ravi   Sharma  Male   30      9876543210      ravi.sharma@example.com
2      Anita  Patel   Female 28      9876543211      anita.patel@example.com
3      Suresh Kumar Male   35      9876543212      suresh.kumar@example.com
Time taken: 0.054 seconds, Fetched: 3 row(s)
```

Transaction

hive> SELECT * FROM Transaction LIMIT 3;

hive> SELECT * FROM Transaction LIMIT 9;

```
hive> SELECT * FROM Transaction LIMIT 3;
OK
101      Payment Received      2023-04-21      123456 201      301
102      Payment Processed      2024-08-30      987654 202      302
103      Payment Completed      2020-06-13      456789 203      303
Time taken: 0.065 seconds, Fetched: 3 row(s)
hive> SELECT * FROM Transaction LIMIT 9;
OK
101      Payment Received      2023-04-21      123456 201      301
102      Payment Processed      2024-08-30      987654 202      302
103      Payment Completed      2020-06-13      456789 203      303
104      Payment Confirmed      2019-10-04      159357 204      304
105      Payment Successful      2024-08-05      753951 205      305
106      Payment Verified      2024-02-26      951357 206      306
107      Payment Approved      2024-04-29      852741 207      307
108      Payment Declined      2021-08-08      147369 208      308
109      Payment Reversed      2016-07-09      963852 209      309
Time taken: 0.061 seconds, Fetched: 9 row(s)
```

3. Display the count of each car type.

```
hive> SELECT cab_type, COUNT(*) AS count
> FROM cab
> GROUP BY cab_type;
```

```
OK
Hatchback      2
Minivan 2
SUV            3
Sedan          3
Time taken: 1.217 seconds, Fetched: 4 row(s)
```

**4. Display the driver details
whose salary is less than 30000.**

```
hive> SELECT *
> FROM Driver
> WHERE driver_sal < 30000;
```

```
hive> SELECT *
> FROM Driver
> WHERE driver_sal < 30000;
OK
1004    Diya Verma      28000
1009    Aditya Joshi    29000
Time taken: 0.21 seconds, Fetched: 2 row(s)
hive>
```

**5. Display the rental details
where the payment is maximum.**

```
hive> SELECT *
> FROM Rental
> WHERE Payment =
(SELECT MAX(Payment) FROM
Rental);
```

```
OK
7      2024-08-07      01:00 PM      Ahmedabad      2100
Time taken: 6.517 seconds, Fetched: 1 row(s)
```

6. Create a partition for customers based on male and female.

```
hive> CREATE VIEW
Male_Customers AS
> SELECT *
> FROM Customer
> WHERE cus_gender = 'Male';
```

```
hive> CREATE VIEW Male_Customers AS
> SELECT *
> FROM Customer
> WHERE cus_gender = 'Male';
OK
Time taken: 0.469 seconds
```

```
hive> select * from male_customers;
OK
1      Ravi      Sharma      Male      30      9876543210      ravi.sharma@example.com
3      Suresh      Kumar      Male      35      9876543212      suresh.kumar@example.com
5      Amit      Singh      Male      40      9876543214      amit.singh@example.com
7      Vikram      Reddy      Male      29      9876543216      vikram.reddy@example.com
9      Ajay      Gupta      Male      38      9876543218      ajay.gupta@example.com
Time taken: 0.171 seconds, Fetched: 5 row(s)
```

```
hive> CREATE VIEW female_Customers AS
> SELECT *
> FROM Customer
> WHERE cus_gender = 'Female';
```

```
hive> CREATE VIEW female_Customers AS
> SELECT *
> FROM Customer
> WHERE cus_gender = 'Female';
OK
Time taken: 0.075 seconds
hive> select * from female_customers;
OK
2      Anita      Patel      Female      28      9876543211      anita.patel@example.com
4      Neha      Verna      Female      26      9876543213      neha.verma@example.com
6      Pooja      Desai      Female      32      9876543215      pooja.desai@example.com
8      Sneha      Mishra      Female      31      9876543217      sneha.mishra@example.com
10     Simran      Kaur      Female      27      9876543219      simran.kaur@example.com
Time taken: 0.062 seconds, Fetched: 5 row(s)
```

```
WHERE cus_gender = 'Female';
```


7. Create three buckets on drivers based on salary.

```
hive> SELECT driver_id, driver_name, driver_sal,  
> CASE  
> WHEN driver_sal < 30000 THEN 'Low'  
> WHEN driver_sal BETWEEN 30000 AND 35000 THEN 'Medium'  
> ELSE 'High'  
> END AS salary_bucket  
> FROM Driver;
```

```
hive> SELECT driver_id, driver_name, driver_sal,  
> CASE  
> WHEN driver_sal < 30000 THEN 'Low'  
> WHEN driver_sal BETWEEN 30000 AND 35000 THEN 'Medium'  
> ELSE 'High'  
> END AS salary_bucket  
> FROM Driver;  
OK  
1001  Aarav Sharma      30000  Medium  
1002  Ananya Gupta       35000  Medium  
1003  Rohan Patel         32000  Medium  
1004  Diya Verma          28000  Low  
1005  Arjun Mehta         40000  High  
1006  Sneha Reddy         45000  High  
1007  Krish Iyer          37000  High  
1008  Pooja Singh         36000  High  
1009  Aditya Joshi        29000  Low  
1010  Neha Nair           31000  Medium  
Time taken: 0.06 seconds, Fetched: 10 row(s)
```

8. Display the transaction details of transactions that happened during the year 2024.

```
hive> SELECT *  
> FROM Transaction  
> WHERE YEAR(tran_date) = 2024;
```

```
hive> SELECT *  
> FROM Transaction  
> WHERE YEAR(tran_date) = 2024;  
OK  
102    Payment Processed      2024-08-30    987654    202    302  
105    Payment Successful     2024-08-05    753951    205    305  
106    Payment Verified       2024-02-26    951357    206    306  
107    Payment Approved       2024-04-29    852741    207    307  
Time taken: 0.078 seconds, Fetched: 4 row(s)
```

9. Display the count of no. of cabs that are available (status='A')

```
hive> SELECT COUNT(*) AS available_cabs  
> FROM cab  
> WHERE cab_status = 'Available';
```

```
OK
7
Time taken: 1.302 seconds, Fetched: 1 row(s)
```

10. Display the average salary of all the drivers.

```
hive> SELECT AVG(driver_sal) AS average_salary
> FROM Driver;
```

```
OK
34300.0
Time taken: 1.573 seconds, Fetched: 1 row(s)
```

5. Data processing using Spark

MovieLens datasets were collected by the GroupLens Research Project at the University of Minnesota. It represents users' reviews of movies.

This data set consists of:

- * 100,000 ratings (1-5) from 943 users on 1682 movies.
 - * Each user has rated at least 20 movies.
 - * Simple demographic info for the users (age, gender, occupation, zip)
- u.data -- The full u data set, 100000 ratings by 943 users on 1682 items.

Each user has rated at least 20 movies. Users and items are numbered consecutively from 1. The data is randomly ordered. This is a tab separated list of user id item id rating timestamp.

The time stamps are unix seconds since 1/1/1970 UTC

u.user -- Demographic information about the users; this is a pipe separated list of user id | age | gender | occupation | zip code

The user ids are the ones used in the u.data data set.

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName("MovieLens Data Analysis").getOrCreate()
sc = spark.sparkContext
```

```
KeyboardInterrupt
>>> from pyspark.sql import SparkSession
>>> spark = SparkSession.builder.appName("MovieLens Data Analysis").getOrCreate()
>>>
>>> sc = spark.sparkContext
>>>
```

1. load the u.data into the new RDD.

#Load u.data into an RDD

```
>> ratings_rdd = sc.textFile("/home/mca/Downloads/u.data")
```

Display the first few records

```
>>> print(ratings_rdd.take(5))
```

```
>>> ratings_rdd = sc.textFile("/home/mca/Downloads/u.data")
>>>
>>> print(ratings_rdd.take(5))
['196\t242\t3\t881250949', '186\t302\t3\t891717742', '22\t377\t1\t878887116', '244\t51\t2\t880606923', '166\t346\t1\t886397596']
>>>
```

```
# Step 4: Parse the RDD into a structured format
parsed_ratings_rdd = ratings_rdd.map(lambda line:line.split('\t'))

>>> print(parsed_ratings_rdd.take(5))
```

```
[('196', '242', '3', '881250949'), ('186', '302', '3', '891717742'),
 ('22', '377', '1', '878887116'), ('244', '51', '2', '880606923'), ('166', '346', '1', '886397596')]
```

2.Change the RDD to a DataFrame.

```
# Convert RDD to DataFrame with appropriate column names
>>> df_ratings = parsed_ratings_rdd.toDF(["user_id", "item_id", "rating", "timestamp"])

ratings_df = spark.createDataFrame(parsed_ratings_rdd, schema=rating_schema)
```

3.Return the schema of this DataFrame.

```
>>> df_ratings.printSchema()
```

```
>>> df_ratings.printSchema()
root
 |-- user_id: string (nullable = true)
 |-- item_id: string (nullable = true)
 |-- rating: string (nullable = true)
 |-- timestamp: string (nullable = true)
```

4.Register the DataFrame as a temp u_data table.

```
>>> df_ratings.createOrReplaceTempView("u_data")
```

5 Display the contents of newly created u_data table

```
>>> df_ratings.show()
```

```
>>> df_ratings.show()
+-----+-----+-----+-----+
|user_id|item_id|rating|timestamp|
+-----+-----+-----+-----+
|196|242|3|881250949|
|186|302|3|891717742|
|22|377|1|878887116|
|244|51|2|880606923|
|166|346|1|886397596|
|298|474|4|884182806|
|115|265|2|881171488|
|253|465|5|891628467|
|305|451|3|886324817|
|6|86|3|883603013|
|62|257|2|879372434|
|286|1014|5|879781125|
|200|222|5|876042340|
|210|40|3|891035994|
|224|29|3|888104457|
|303|785|3|879485318|
|122|387|5|879270459|
|194|274|2|879539794|
|291|1042|4|874834944|
|234|1184|2|892079237|
+-----+-----+-----+-----+
only showing top 20 rows
```

6. Show the numbers of items reviewed by each user in the newly created u_data table.

```
>>> review_counts_by_user = spark.sql("""
...     SELECT user_id, COUNT(item_id) AS num_items_reviewed
...     FROM u_data
...     GROUP BY user_id
...     ORDER BY num_items_reviewed DESC """)
```

```
>>> review_counts_by_user.show()
```

```
>>> review_counts_by_user.show()
+-----+-----+
|user_id|num_items_reviewed|
+-----+-----+
|405|737|
|655|685|
|13|636|
|450|540|
|276|518|
|416|493|
|537|490|
|303|484|
|234|480|
|393|448|
|181|435|
|279|434|
|429|414|
|846|405|
|7|403|
|94|400|
|682|399|
|308|397|
|92|388|
|293|388|
+-----+-----+
only showing top 20 rows
```

7. Show the numbers of users reviewed each item in the newly created u_data table

```
>>> review_counts_by_item = spark.sql("""
...     SELECT item_id, COUNT(user_id) AS num_users_reviewed
...     FROM u_data
...     GROUP BY item_id
...     ORDER BY num_users_reviewed DESC """)
```

```
>>> review_counts_by_item.show()
```

```
>>> review_counts_by_item.show()
+-----+-----+
|item_id|num_users_reviewed|
+-----+-----+
| 50|583|
|258|509|
|100|508|
|181|507|
|294|485|
|286|481|
|288|478|
| 1|452|
|300|431|
|121|429|
|174|420|
|127|413|
| 56|394|
| 7|392|
| 98|390|
|237|384|
|117|378|
|172|367|
|222|365|
|313|350|
+-----+-----+
only showing top 20 rows
```

8. Load the u.user into a new RDD.

```
>>> users_rdd = sc.textFile("/home/mca/Downloads/u.user")
```

```
>>> print(users_rdd.take(5))
```

```
>>> print(users_rdd.take(5))
'1|24|M|technician|85711', '2|53|F|other|94043',
>>>
'3|23|M|writer|32067', '4|24|M|technician|43537', '5|33|F|other|15213']
```

9. Change the RDD to a Dataframe.

```
>>> parsed_users_rdd = users_rdd.map(lambda line: line.split('|'))
```

```
>>> df_users = parsed_users_rdd.toDF(["user_id", "age", "gender", "occupation",
"zip_code"])
```

```
>>> df_users.printSchema()
```

```
>>> df_users.printSchema()
root
 |-- user_id: string (nullable = true)
 |-- age: string (nullable = true)
 |-- gender: string (nullable = true)
 |-- occupation: string (nullable = true)
 |-- zip_code: string (nullable = true)

>>> df_users.printSchema()
root
 |-- user_id: string (nullable = true)
 |-- age: string (nullable = true)
 |-- gender: string (nullable = true)
 |-- occupation: string (nullable = true)
 |-- zip_code: string (nullable = true)
```

10 Register the DataFrame as a temp u_user table.

```
>>> df_users.createOrReplaceTempView("u_user")
```

11. Display the contents of newly created user table

```
>>> # Display the contents of the u_user table
```

```
>>> df_users.show()
```

```
>>> df_users.show()
+-----+-----+-----+-----+-----+
|user_id|age|gender|  occupation|zip_code|
+-----+-----+-----+-----+-----+
|    1| 24|   M|   technician|  85711|
|    2| 53|   F|         other|  94043|
|    3| 23|   M|        writer|  32067|
|    4| 24|   M|   technician|  43537|
|    5| 33|   F|         other|  15213|
|    6| 42|   M|   executive|  98101|
|    7| 57|   M| administrator|  91344|
|    8| 36|   M| administrator|  05201|
|    9| 29|   M|        student|  01002|
|   10| 53|   M|        lawyer|  90703|
|   11| 39|   F|         other|  30329|
|   12| 28|   F|         other|  06405|
|   13| 47|   M|       educator|  29206|
|   14| 45|   M|    scientist|  55106|
|   15| 49|   F|       educator|  97301|
|   16| 21|   M| entertainment|  10309|
|   17| 30|   M|   programmer|  06355|
|   18| 35|   F|         other|  37212|
|   19| 40|   M|    librarian|  02138|
|   20| 42|   F|    homemaker|  95660|
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

12.Count the number of user in the u_user table gender wise

```
>>> gender_counts = spark.sql("""
...     SELECT gender, COUNT(user_id) AS num_users
...     FROM u_user
...     GROUP BY gender
... """)
```

```
>>> gender_counts.show()
```

```
... )
>>> gender_counts.show()
+-----+-----+
|gender|num_users|
+-----+-----+
|      F|      273|
|      M|      670|
+-----+-----+
```

13. Join u_data table and u_user tables based on userid

```
>>> df_ratings.createOrReplaceTempView("u_data")
```

```
>>> df_users.createOrReplaceTempView("u_user")
```

```
>>> joined_df = spark.sql("""
...     SELECT d.user_id, d.item_id, d.rating, d.timestamp, u.age, u.gender, u.occupation,
...     u.zip_code
...     FROM u_data d
...     JOIN u_user u
...     ON d.user_id = u.user_id
... """)
>>> joined_df.show()
```

```
>>> joined_df.show()
+-----+-----+-----+-----+-----+-----+-----+
|user_id|item_id|rating|timestamp|age|gender|occupation|zip_code|
+-----+-----+-----+-----+-----+-----+-----+
| 296| 705| 5|884197193| 43| F|administrator| 16803|
| 296| 508| 5|884196584| 43| F|administrator| 16803|
| 296| 20| 5|884196921| 43| F|administrator| 16803|
| 296| 228| 4|884197264| 43| F|administrator| 16803|
| 296| 222| 5|884196640| 43| F|administrator| 16803|
| 296| 429| 5|884197330| 43| F|administrator| 16803|
| 296| 855| 5|884197352| 43| F|administrator| 16803|
| 296| 248| 5|884196765| 43| F|administrator| 16803|
| 296| 258| 5|884196469| 43| F|administrator| 16803|
| 296| 242| 4|884196057| 43| F|administrator| 16803|
| 296| 48| 5|884197091| 43| F|administrator| 16803|
| 296| 286| 5|884196209| 43| F|administrator| 16803|
| 296| 272| 5|884198772| 43| F|administrator| 16803|
| 296| 510| 5|884197264| 43| F|administrator| 16803|
| 296| 275| 4|884196555| 43| F|administrator| 16803|
| 296| 427| 5|884198772| 43| F|administrator| 16803|
| 296| 83| 5|884199624| 43| F|administrator| 16803|
| 296| 961| 5|884197287| 43| F|administrator| 16803|
| 296| 544| 4|884196938| 43| F|administrator| 16803|
| 296| 32| 4|884197131| 43| F|administrator| 16803|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```


6. Programming in Cassandra

1. Create KeySpace “Students”

```
>> CREATE KEYSPACE Students WITH replication = {'class': 'SimpleStrategy',
'replication_factor': 1};
```

2. Describe the existing Keyspaces

```
>> DESCRIBE KEYSPACES;
```

```
cqlsh> DESCRIBE KEYSPACES;

system_virtual_schema  system_schema  system_views  system_distributed
students                system_auth    system         system_traces
```

3. DESCRIBE KEYSPACE Students;

```
cqlsh> SELECT * FROM students.schema_keyspaces;
```

```
cqlsh> SELECT * FROM system_schema.keyspaces;
keyspace_name | durable_writes | replication
-----
system_auth   | True           | {'class': 'org.apache.cassandra.locator.SimpleStrategy', 'replication_factor': '1'}
system_schema | True           | {'class': 'org.apache.cassandra.locator.LocalStrategy'}
system_distributed | True         | {'class': 'org.apache.cassandra.locator.SimpleStrategy', 'replication_factor': '3'}
system        | True           | {'class': 'org.apache.cassandra.locator.LocalStrategy'}
system_traces | True           | {'class': 'org.apache.cassandra.locator.SimpleStrategy', 'replication_factor': '2'}
students      | True           | {'class': 'org.apache.cassandra.locator.SimpleStrategy', 'replication_factor': '1'}
```

4. Use the keyspace “Students”

```
cqlsh:students> USE Students;
```

5. Student details:

Create table (column family) by name Student_Info(RollNo int, StudName text, DateOfJoining timestamp, PrevSemPercentage double), RollNo is primary key

```
cqlsh:students> CREATE TABLE Student_Info (
    RollNo int PRIMARY KEY,
    StudName text,
    DateOfJoining timestamp,
    PrevSemPercentage double
);
```

```
cqlsh> USE Students;
cqlsh:students> CREATE TABLE Student_Info (
    RollNo int PRIMARY KEY,
    StudName text,
    DateOfJoining timestamp,
    PrevSemPercentage double
);
```

6.Book_Borrowed

Create table (column family) by name **Library_Book** (CountValue counter, BookName varchar, RollNo int, StudName varchar), PRIMARY KEY is (book_name,stud_name));

```
cqlsh:students> CREATE TABLE Library_Book_Counter (
    BookName varchar,
    StudName varchar,
    CountValue counter,
    PRIMARY KEY (BookName, StudName)
);
```

```
cqlsh:students> CREATE TABLE Library_Book (
    BookName varchar,
    RollNo int,
    StudName varchar,
    PRIMARY KEY (BookName, StudName)
);
```

```
cqlsh:students>
cqlsh:students> CREATE TABLE Library_Book_Counter (
    BookName varchar,
    StudName varchar,
    CountValue counter,
    PRIMARY KEY (BookName, StudName)
);
cqlsh:students> CREATE TABLE Library_Book (
    BookName varchar,
    RollNo int,
    StudName varchar,
    PRIMARY KEY (BookName, StudName)
);
cqlsh:students>
```

7.Lookup the names of all tables in the current keyspace

cqlsh:students> DESCRIBE TABLES;

```
cqlsh:students> DESCRIBE TABLES;

library_book_counter  student_info  library_book
```

8.Describe the table information

cqlsh:students> DESCRIBE TABLE Student_Info;

```
CREATE TABLE students.student_info (
  rollno int PRIMARY KEY,
  dateofjoining timestamp,
  prevsempersent double,
  studname text
) WITH additional_write_policy = '99p'
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
AND crc_check_chance = 1.0
AND default_time_to_live = 0
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair = 'BLOCKING'
AND speculative_retry = '99p';
```

cqlsh:students> DESCRIBE TABLE Library_Book;

```
cqlsh:students> DESCRIBE TABLE Library_Book;

CREATE TABLE students.library_book (
  bookname text,
  studname text,
  rollno int,
  PRIMARY KEY (bookname, studname)
) WITH CLUSTERING ORDER BY (studname ASC)
AND additional_write_policy = '99p'
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
AND crc_check_chance = 1.0
AND default_time_to_live = 0
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair = 'BLOCKING'
AND speculative_retry = '99p';
```

9.CRUD**Insert at least 5 rows for Student_Info**

cqlsh:students>

```
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage)
VALUES (101, 'Aarav', '2024-01-15', 85.5);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage)
VALUES (102, 'Vivaan', '2024-01-16', 90.0);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage)
VALUES (103, 'Aditya', '2024-01-17', 78.0);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage)
VALUES (104, 'Vihaan', '2024-01-18', 88.0);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage)
VALUES (105, 'Reyansh', '2024-01-19', 92.0);
```

```
cqlsh:students> INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage) VALUES (101, 'Aarav', '2024-01-15', 85.5);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage) VALUES (102, 'Vivaan', '2024-01-16', 90.0);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage) VALUES (103, 'Aditya', '2024-01-17', 78.0);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage) VALUES (104, 'Vihaan', '2024-01-18', 88.0);
INSERT INTO Student_Info (RollNo, StudName, DateOfJoining, PrevSemPercentage) VALUES (105, 'Reyansh', '2024-01-19', 92.0);
```

10.View data from the table “Students_Info”

cqlsh:students> SELECT * FROM Student_Info;

```
cqlsh:students> SELECT * FROM Student_Info;
```

rollno	dateofjoining	prevsempercentage	studname
5	2024-01-18 18:30:00.000000+0000	92	Reyansh
1	2024-01-14 18:30:00.000000+0000	85.5	Aarav
2	2024-01-15 18:30:00.000000+0000	90	Vivaan
4	2024-01-17 18:30:00.000000+0000	88	Vihaan
3	2024-01-16 18:30:00.000000+0000	78	Aditya

(5 rows)

11.View data from the table “Students_Info” where RollNo column either has a value 1 or 2 or 3

cqlsh:students> SELECT * FROM Student_Info WHERE RollNo IN (1, 2, 3);

```
cqlsh:students> SELECT * FROM Student_Info WHERE RollNo IN (1, 2, 3);
```

rollno	dateofjoining	prevsempercentage	studname
1	2024-01-14 18:30:00.000000+0000	85.5	Aarav
2	2024-01-15 18:30:00.000000+0000	90	Vivaan
3	2024-01-16 18:30:00.000000+0000	78	Aditya

(3 rows)

12. To execute a non primary key**Create an INDEX on the Column StudName**

```
cqlsh:students> CREATE INDEX ON Student_Info (StudName);
```

13. Execute the query based on the INDEXED Column:**Display students details for a specific student name.**

```
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aarav';
```

```
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aditya';
```

```
cqlsh:students> CREATE INDEX ON Student_Info (StudName);
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aarav';

rollno | dateofjoining | prevsempercentage | studname
-----+-----+-----+-----
1 | 2024-01-14 18:30:00.000000+0000 | 85.5 | Aarav
(1 rows)
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aditya';

rollno | dateofjoining | prevsempercentage | studname
-----+-----+-----+-----
3 | 2024-01-16 18:30:00.000000+0000 | 78 | Aditya
(1 rows)
```

14. Specify the number of rows to display in the output

```
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aditya' LIMIT 1;
```

```
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aditya' LIMIT 1;

rollno | dateofjoining | prevsempercentage | studname
-----+-----+-----+-----
3 | 2024-01-16 18:30:00.000000+0000 | 78 | Aditya
(1 rows)
```

```
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aditya' LIMIT 2;
```

```
(1 rows)
cqlsh:students> SELECT * FROM Student_Info WHERE StudName = 'Aditya' LIMIT 2;

rollno | dateofjoining | prevsempercentage | studname
-----+-----+-----+-----
7 | 2024-01-13 18:30:00.000000+0000 | 58 | Aditya
6 | 2024-01-13 18:30:00.000000+0000 | 98 | Aditya
(2 rows)
```

15. Alias for Column:

Display RollNo as "USN"

```
cqlsh:students> SELECT RollNo AS "USN", StudName, DateOfJoining, PrevSemPercentage
FROM Student_Info;
```

```
cqlsh:students> SELECT RollNo AS "USN", StudName, DateOfJoining, PrevSemPercentage FROM Student_Info;
```

USN	studname	dateofjoining	prevsempercentage
5	Reyansh	2024-01-18 18:30:00.000000+0000	92
1	Aarav	2024-01-14 18:30:00.000000+0000	85.5
2	Vivaan	2024-01-15 18:30:00.000000+0000	90
4	Vihaan	2024-01-17 18:30:00.000000+0000	88
7	Aditya	2024-01-13 18:30:00.000000+0000	58
6	Aditya	2024-01-13 18:30:00.000000+0000	98
3	Aditya	2024-01-13 18:30:00.000000+0000	98

(7 rows)

16.UPDATE the student name with last name for a specific RollNo

```
cqlsh:students> UPDATE Student_Info SET StudName = 'Aarav Smitha' WHERE RollNo = 1;
```

```
cqlsh:students> SELECT * FROM Student_Info where rollno=1;
```

rollno	dateofjoining	prevsempercentage	studname
1	2024-01-14 18:30:00.000000+0000	85.5	Aarav

(1 rows)

```
cqlsh:students> UPDATE Student_Info SET StudName = 'Aarav Smitha' WHERE RollNo = 1;
cqlsh:students> SELECT * FROM Student_Info where rollno=1;
```

rollno	dateofjoining	prevsempercentage	studname
1	2024-01-14 18:30:00.000000+0000	85.5	Aarav Smitha

Change the RollNo to 10 for a existing RollNo with value 1.

```
cqlsh:students> INSERT INTO Student_Info (RollNo, StudName, DateOfJoining,
PrevSemPercentage)
VALUES (10, 'Aarav Smitha', '2024-01-15', 85.5);
```

```
cqlsh:students> DELETE FROM Student_Info WHERE RollNo = 1;
```

```
cqlsh:students> SELECT * FROM Student_Info where rollno=10;
```

rollno	dateofjoining	prevsempercentage	studname
10	2024-01-14 18:30:00.000000+0000	85.5	Aarav Smitha

17.DELETE PrevSemPercent for student with RollNo=2;

cqlsh:students> UPDATE Student_Info SET PrevSemPercentage = null WHERE RollNo = 2;

```
cqlsh:students> SELECT * FROM Student_Info;
```

rollno	dateofjoining	prevsempercentage	studname
5	2024-01-18 18:30:00.000000+0000	92	Reyansh
10	2024-01-14 18:30:00.000000+0000	85.5	Aarav Smitha
2	2024-01-15 18:30:00.000000+0000	null	Vivaan
4	2024-01-17 18:30:00.000000+0000	88	Vihaan
7	2024-01-13 18:30:00.000000+0000	58	Aditya
6	2024-01-13 18:30:00.000000+0000	98	Aditya
3	2024-01-13 18:30:00.000000+0000	98	Aditya

18.Delete a Row FROM student_info WHERE RollNo is 3;

cqlsh:students> DELETE FROM Student_Info WHERE RollNo = 3;

```
cqlsh:students> SELECT * FROM Student_Info;
```

rollno	dateofjoining	prevsempercentage	studname
5	2024-01-18 18:30:00.000000+0000	92	Reyansh
10	2024-01-14 18:30:00.000000+0000	85.5	Aarav Smitha
2	2024-01-15 18:30:00.000000+0000	null	Vivaan
4	2024-01-17 18:30:00.000000+0000	88	Vihaan
7	2024-01-13 18:30:00.000000+0000	58	Aditya
6	2024-01-13 18:30:00.000000+0000	98	Aditya

19.Set Collection

A column of type set consists of unordered unique values. However, when the column is queried, it returns the values in sorted order. For example, for text values, it sorts in alphabetical order.

cqlsh:students> ALTER TABLE Student_Info ADD hobbies SET<text>;

```
cqlsh:students> describe student_info;
```

```
CREATE TABLE students.student_info (
  rollno int PRIMARY KEY,
  dateofjoining timestamp,
  hobbies set<text>,
  prevsempercentage double,
  studname text
) WITH additional_write_policy = '99p'
```

20.List Collection

When the order of elements matter, one should go for a list collection.

Alter the StudentsInfo table to add language as a list of text

```
cqlsh:students> ALTER TABLE Student_Info ADD languages LIST<text>;
```

```
CREATE TABLE students.student_info (
  rollno int PRIMARY KEY,
  dateofjoining timestamp,
  hobbies set<text>,
  languages list<text>,
  prevsempercentage double,
  studname text
) WITH additional_write_policy = '99p'
AND bloom_filter_chance = 0.01
```

21.Update the values for hobbies column (Music Cricket) and language column (Kannada, Hindi,English) for RollNo with value 10 and display the student-info

```
cqlsh:students> UPDATE Student_Info SET hobbies = {'Music', 'Cricket'}, languages = ['Kannada', 'Hindi', 'English'] WHERE RollNo = 10;
```

```
cqlsh:students> SELECT * FROM Student_Info where rollno=10;
```

rollno	dateofjoining	hobbies	languages	prevsempercentage	studname
10	2024-01-14 18:30:00.000000+0000	['Cricket', 'Music']	['Kannada', 'Hindi', 'English']	85.5	Aarav Smitha

(1 rows)

22.Remove Hindi from the language list for RollNo 10 and display the student info

```
cqlsh:students> UPDATE Student_Info SET languages = languages - ['Hindi'] WHERE RollNo = 10;
```

```
cqlsh:students> SELECT * FROM Student_Info where rollno=10;
```

rollno	dateofjoining	hobbies	languages	prevsempercentage	studname
10	2024-01-14 18:30:00.000000+0000	['Cricket', 'Music']	['Kannada', 'English']	85.5	Aarav Smitha

(1 rows)

23.USING A COUNTER

A counter is a special column that is changed in increments. For example, we may need a counter column to count the number of times a particular book is issued from the library by the student.

```
cqlsh:students> UPDATE Library_Book_Counter SET CountValue = CountValue +  
1 WHERE BookName = 'Big Data Analytics' AND StudName = 'Ram';
```

```
cqlsh:students> SELECT * FROM Library_Book_Counter WHERE BookName = 'Big Data Analytics' AND StudName = 'Ram';  
  
bookname      | studname | countvalue  
-----  
Big Data Analytics | Ram | 1  
(1 rows)  
  
cqlsh:students> SELECT * FROM Library_Book_Counter WHERE BookName = 'Big Data Analytics' AND StudName = 'Ram';  
  
bookname      | studname | countvalue  
-----  
Big Data Analytics | Ram | 2  
(1 rows)
```